

A governed AI workflow for public-record review

Prepared by Tim Sullivan | August 11, 2026 | Conceptual Creatio mapping; no Creatio deployment is claimed.

OBJECTIVE: Reduce the manual effort required to assemble address-based public records while preserving exact sources, unresolved fields, and human control over consequential output.

Current process

A property or community review can require separate federal, state, county, regional, and municipal searches. Sources use different address formats, update schedules, query interfaces, and meanings. An unavailable record cannot safely be converted into “none exists,” and a screening result cannot become a legal, engineering, insurance, environmental, lending, or regulatory determination.

Stakeholders and decisions

- Municipal or regional staff: locate records, create an inventory, and route unresolved items.
- Environmental and engineering consultants: assemble evidence before fieldwork or professional review.
- Property professionals: identify which official records require attention before a transaction or project.
- Review owner: decide whether the evidence is sufficient to release, escalate, or leave unresolved.

Use-case priority

Use case	Business value	Feasibility / risk	Current state
One-address, source-cited review	High: replaces repeated portal searches	High feasibility; bounded source claims	Implemented broadly
Unresolved-item queue	High: prevents silent gaps and directs follow-up	High feasibility; requires ownership and due dates	Implemented as product state
Community address filtering	High for inventories and prioritization	Feasible where a denominator and layer are available	Implemented in bounded workflows
Multi-town coalition view	High for regional buyers	Requires generalized jurisdiction and denominator model	Roadmap; not generalized

Priority rule: automate repetitive, evidence-bounded steps first; retain human review where an error changes an external decision.

Proposed discovery questions

1. Which decision is the review supporting, and which records are authoritative for that decision?
2. What is the full denominator: one address, every parcel in one jurisdiction, or a multi-town region?
3. What output requires human approval, and who owns unresolved items?
4. Which measures establish value: cycle time, exact-source coverage, error rate, unresolved closure, or cost per completed review?

Agents accelerate the work; evidence controls the answer

The model is used where judgment and synthesis help. Deterministic services retain responsibility for address matching and compliance-adjacent result fields.



Proposed Creatio implementation

- Record model: a case, property, or portfolio record stores the request, jurisdiction, review stage, returned findings, sources, owners, approvals, and unresolved work.
- Business Studio: orchestrates intake, connector calls, review tasks, approval steps, service levels, and exception paths.
- AI Studio: agents triage candidate sources, summarize conflicts, prepare follow-up tasks, and draft review notes inside a constrained schema.
- Integration layer: REST and MCP tools separate retrieval from verification; credentials and permissions remain outside model context.
- Release control: a human approval step controls external delivery whenever a result is incomplete, conflicting, or consequential.
- Operations: dashboards report cycle time, exact-source coverage, unresolved aging, exception rate, and cost per completed review.

Guardrail contract

Control	Enforcement	Failure behavior
Evidence	Official source, URL, checked date, scope and adapter version	Hold result or mark unresolved
Entity resolution	Normalized address plus positive controls on known records	No match is not treated as no record
Language	Allowed statements derived from source-support contract	Use “Not confirmed”; prohibit determination language
Permissions	Role-based access; secrets remain outside prompts	Deny tool call and create review task
Release	Computed, reviewed, customer-visible and public states remain separate	No external release without required approval
Monitoring	Golden-set regression checks and exception sampling	Rollback or quarantine affected source lane

DESIGN DECISION: The customer-facing answer path remains reproducible. Agents can propose, summarize and route work; they do not convert absent evidence into a consequential property determination.

A bounded deployment with evidence before scale

Start with one jurisdiction, one source family, one review owner, and a golden set of known records.

Four-stage pilot

- 1. Discovery** Map current steps, decisions, systems, owners, permissions, volumes, baseline time and failure modes.
- 2. Prototype** Configure the record model and one workflow; connect one source family; define agent tools, approvals and the golden test set.
- 3. Controlled pilot** Run live work with human review; record exceptions, unresolved reasons, review time and user feedback.
- 4. Scale decision** Expand only if source accuracy, workflow adoption, cycle-time reduction and exception handling meet agreed thresholds.

Scorecard

Measure	Definition	Pilot target
Median cycle time	Intake to reviewed output	Set after baseline; no invented target
Exact-source coverage	Released findings with reviewable supporting URL	Agree by source family
Golden-set accuracy	Correct supported result or correct abstention	Set before live pilot
Unresolved aging	Time until closure, escalation or documented stop	Service level set by owner
Human-review rate	Share requiring intervention and reason	Expected to decline only without accuracy loss
Cost per completed review	Integration, compute and reviewer time per accepted output	Compare with manual baseline

Working evidence available today

- Lasting Ground public validation examples: 18 passing tests at 93% line coverage.
- Agentic Engineering: a deterministic-plus-rubric evaluation gate, drift check, Agent Skill, and typed MCP retrieval and verification tools.
- GroundTruth-Geo: 33 government-cited benchmark items across seven states with a deterministic grader and MCP interface.
- SquadBrain: a shipped iOS product; public TypeScript logic examples pass 25 tests at 98.3% line coverage.

Current product limits

Multi-town filtering is not generalized. State and country depth is uneven. Some rows identify an indexed official source rather than a record-level deep link. Lasting Ground has not yet had customers or received payments; this is a product-and-systems example, not a customer case study.

Review links: [Portfolio](#) | [Lasting Ground examples](#) | [Agent workflow examples](#) | [Creatio role](#)